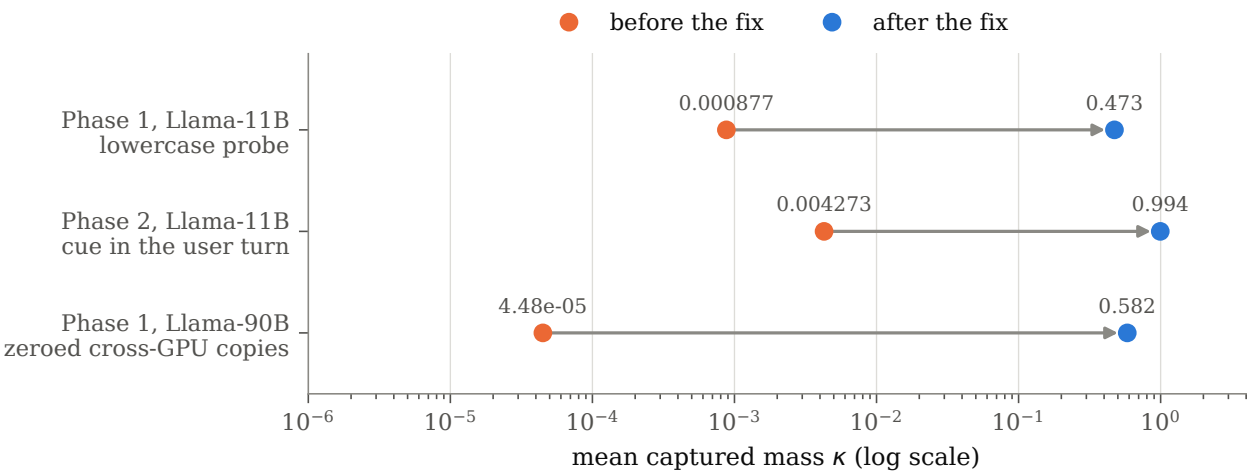
before the fix
after the fix
Phase 1, Llama-11B
lowercase probe
0.000877
0.473
Phase 2, Llama-11B
cue in the user turn
0.004273
0.994
Phase 1, Llama-90B
zeroed cross-GPU copies
4.48e-05
0.582
$10^{-6}$
$10^{-5}$
$10^{-4}$
$10^{-3}$
$10^{-2}$
$10^{-1}$
$10^{0}$
mean captured mass $\kappa$ (log scale)